

TEJA SAI SRINIVAS

+1 940-597-2906 | tejasaisrin22@gmail.com | [LinkedIn](#) | [GitHub](#)

SUMMARY

Data Science Master's student (4.0/4.0 GPA) specializing in AI/ML systems, with hands-on experience building LLM-powered applications, RAG pipelines, and production-grade APIs using Python, FastAPI, and LangChain. Proven track record deploying multi-agent AI systems and full-stack platforms from design to production. Passionate about applying AI infrastructure to high-impact domains like healthcare and clinical intelligence.

EDUCATION

University of North Texas

Aug 2024 – May 2026 (Exp)

Master of Science, Data Science (STEM), GPA: 4.0

Shiv Nadar Institute of Eminence

Aug 2020 - May 2024

Bachelor's of Technology, Computer Science and Engineering

Specialization in Data Science with a Minor in Mathematics

EXPERIENCE

Cargo Exchange India Private Limited

Jun 2023 - Aug 2023

Data Analyst Intern

Hyderabad, India

Project: Freight Index System

- Built automated ETL pipelines in Python processing 50K+ daily freight records, reducing manual reporting time by ~40% and surfacing cost-reduction opportunities across 12 logistics routes.
- Developed Power BI dashboards tracking 8+ operational KPIs (route efficiency, capacity utilization, delivery performance), enabling data-driven decisions adopted by senior leadership.
- Analyzed pricing dynamics and demand patterns across the freight vertical to inform revenue optimization strategies; improvements contributed to measurable gains in forecasting accuracy and operational consistency.

PROJECTS

Multi-Agent System for Autonomous Cyber Defense — FastAPI, Kafka, Neo4j, FAISS

[Github](#)

- Architected an agentic AI system that autonomously detects and classifies cyber threats in real-time, achieving 92%+ accuracy across a 10,000-event benchmark dataset.
- Implemented a RAG pipeline using FAISS vector search and a Neo4j knowledge graph, reducing LLM hallucination rate by ~35% versus a vanilla prompting baseline across a 500-case evaluation set.
- Built a microservices architecture using FastAPI and Kafka for event-driven agent communication, achieving avg. response latency under 800ms for real-time security incident triage.
- Deployed local LLMs (Ollama) for on-premise sensitive data handling, reducing avg. inference time from 4.2s to 1.1s through quantization and request-batching optimizations.

Pathfinder (Career Matching Platform) - React, FastAPI, PostgreSQL, TypeScript

[Github](#)

- Developed a full-stack career matching system implementing a custom weighted Jaccard Similarity algorithm against O*NET occupational data, achieving 87% candidate-role match accuracy across 1,200+ profiles.
- Engineered a real-time application lifecycle management system with TanStack Query and TypeScript, featuring role-based dashboards, automated interview scheduling, and instant status notifications.
- Designed Recharts analytics dashboards visualizing candidate placement rates and skill gap analysis for academic advisors, reducing manual reporting effort by an estimated 60%.

Crop Health Monitoring System - Python, Sentinel-2, XGBoost, LSTM, Streamlit, DTW

[Github](#)

- Engineered an end-to-end remote sensing pipeline processing Sentinel-2 satellite imagery to automate NDVI/EVI feature extraction for 2,400+ crop patches across three growing seasons.
- Developed predictive models achieving 0.08 Test RMSE in yield forecasting (XGBoost) and 85% accuracy in stress classification (LSTM); deployed via interactive Streamlit dashboard for real-time stakeholder monitoring.
- Implemented Dynamic Time Warping (DTW) for time-series clustering, identifying 6 distinct crop growth signatures used to guide regional irrigation planning.

Lumen Learning Management System (LLMS) - PostgreSQL, Node.js, Express.js, Docker, SQL, JWT

[Github](#)

- Developed a scalable RESTful API handling JWT authentication and high-volume concurrent requests for a learning platform serving 1,000+ users; achieved avg. query latency < 200ms under 500 concurrent users.
- Designed a normalized PostgreSQL schema with optimized indexing and query plans, eliminating N+1 query patterns and reducing database load by 45%.

- Containerized services with Docker ensuring dev/production parity; implemented RBAC and achieved 95% unit test coverage across all authentication and data endpoints.

TECHNICAL SKILLS

- Languages:** Java, Python (Advanced), SQL, JavaScript/TypeScript, C
- AI & GenAI:** LLMs (Llama, Ollama, Hugging Face), RAG Pipelines, LangChain, Prompt Engineering, FAISS, Agent Architectures
- Cloud & Infrastructure:** GCP (Vertex AI, BigQuery), AWS (S3, Lambda), Docker, Kafka, FastAPI
- ML & Analytics:** PyTorch, XGBoost, Scikit-learn, MLflow, Time Series Forecasting, Statistical Inference, Model Evaluation & Monitoring
- Databases:** PostgreSQL, Neo4j, Pinecone
- Tools:** Git, REST APIs, CI/CD, Unit Testing, Agile/Scrum

CERTIFICATIONS

- Oracle Cloud Infrastructure 2025 Data Science Professional:**

[*View Badge*](#)

- UNT Certified Microcredentials:**

AI Fundamentals Part 1: Demystifying AI,
AI Fundamentals Part 2: How You Can Use AI